

Parallel Signal Processing: Moving-Average Noise Reduction

Programming Assignment 2 (PA2) Comprehensive Final Report

Integrating Sequential, OpenMP, CUDA Baseline, CUDA Optimized, Basic MPI, and Optimized MPI Paradigms with Hardware Energy Profiling and Nsight Systems Diagnostics

Front Matter: Assignment Information

Metadata Item	Project Details / Verified Records
Course	CSS311: Parallel & Distributed Computing
Assignment	Programming Assignment 2 (PA2 Comprehensive Final Report)
Application Theme	Signal Processing
Assigned Problem	1D Moving-Average Noise Reduction Filter Across Six Parallel Paradigms
Group Number	16
Group Leader	Saksham Singh — 2024BCS0070
Group Member 2	Daksh Singh — 2024BCS0042
Group Member 3	Anmol Pipara — 2024BCS0014
Demonstration Link	https://drive.google.com/drive/folders/1i3eLAEA6jXPNDnx8pnbJffigIk-LwEPE?usp=sharing

1. Problem Statement & Mathematical Foundation

Digital signal acquisition systems—such as acoustic transceivers, biomedical monitors (e.g., electrocardiograms and electromyograms), environmental sensor networks, seismic sensors, and radar receivers—continuously capture discrete-time physical signals. During acquisition and physical transmission, these signals are corrupted by high-frequency disturbances, including thermal Johnson–Nyquist noise, electromagnetic interference, and analog-to-digital quantization artifacts.

A fundamental digital signal processing technique for conditioning 1D telemetry is the **Moving-Average (MA) Filter**. The moving-average filter is an unweighted, linear, finite impulse response (FIR) low-pass filter. For each sample in an input discrete-time signal x of length N , the filter computes the unweighted arithmetic mean of W contiguous samples symmetrically centered around the target index:

$$y[i] = \frac{1}{W} \sum_{j=-k}^k x[\text{boundary}(i+j)] \quad \text{for } i \in [0, N-1]$$

where:

- N is the total number of discrete samples in the 1D signal.
- W is an odd positive integer denoting the filter window size ($W = 2k + 1$).
- $k = \frac{W-1}{2}$ represents the filter half-width (radius).
- $\text{boundary}(p)$ maps an arbitrary offset index $p = i + j$ into the valid discrete signal domain $[0, N-1]$.

Boundary Handling: Edge-Clamping (Replication)

At the signal extremities ($i < k$ or $i \geq N - k$), the filter window extends beyond the boundaries of the input array. In this project, all six parallel implementations execute identical **edge-clamping boundary handling** (nearest-element replication):

$$\text{boundary}(p) = \min(\max(p, 0), N-1)$$

Under this policy:

- When $p < 0$, the signal value is clamped to $x[0]$.
- When $p \geq N$, the signal value is clamped to $x[N-1]$.
- When $0 \leq p \leq N-1$, the index remains unchanged as $x[p]$.

Computational Properties of Edge Clamping

- Constant Computational Weight:** Every output element $y[i]$ performs the exact same arithmetic workload: W memory fetches, $W-1$ floating-point additions, and 1 multiplication by the precomputed reciprocal $\frac{1}{W}$.
- Branch Convergence in SIMT/SIMD:** Unlike zero-padding or dynamic window truncation, edge clamping avoids runtime conditional divergence across adjacent threads and vector lanes, guaranteeing uniform instruction execution.
- Physical Signal Continuity:** Clamping eliminates artificial high-frequency transient steps that zero-padding introduces at signal edges, maintaining baseline continuity.

2. Executive Summary & Objectives of PA2

Programming Assignment 1 (PA1) investigated single-node shared-memory and GPU acceleration paradigms:


- Sequential CPU Baseline (single-threaded C++)
- OpenMP Multi-Core Parallelism (multi-threaded CPU)
- CUDA Baseline (global memory many-core GPU)
- CUDA Optimized (on-chip shared-memory tiling with cooperative halo caching)

Programming Assignment 2 (PA2) extends this research into distributed-memory message-passing and deep hardware profiling:

- Basic MPI:** 1D domain decomposition with blocking point-to-point halo exchange (`MPI_Sendrecv`).
- Optimized MPI:** 1D domain decomposition with non-blocking halo exchange (`MPI_Isend` / `MPI_Irecv`) and fine-grained computation/communication overlap (splitting execution into interior stencil computation and boundary stencil computation).
- Hardware Energy Telemetry:** Empirical measurement of CPU package energy via Intel Running Average Power Limit (RAPL) performance counters and GPU board energy via NVIDIA Management Library (NVML) APIs.
- CUDA Profiling via NVIDIA Nsight Systems:** Quantitative analysis of kernel execution durations, memory copy overheads, and hardware resource utilization.
- Six-Way Cross-Paradigm Benchmark:** Direct empirical comparison of all six implementations across Small ($N = 10,000$, $W = 15$), Medium ($N = 100,000$, $W = 31$), and Large ($N = 1,000,000$, $W = 63$) workloads.

All six implementations execute the exact same algorithmic task, consume identical input datasets, apply the identical mathematical formulation and edge-clamping boundary logic, and are rigorously verified for numerical equivalence.

3. The Six Computing Paradigms & Implementations



3.1 Sequential CPU Baseline

The sequential implementation serves as the foundational algorithmic reference. It iterates sequentially over $i \in [0, N - 1]$, accumulating W window elements into a double-precision accumulator and scaling by $\frac{1.0}{W}$.

```
void moving_average_seq(const double* input, double* output, int N, int window_size) {
    int k = (window_size - 1) / 2;
    double inv_w = 1.0 / window_size;
    for (int i = 0; i < N; ++i) {
        double sum = 0.0;
        for (int j = -k; j <= k; ++j) {
            int idx = i + j;
            if (idx < 0) idx = 0;
            else if (idx >= N) idx = N - 1;
            sum += input[idx];
        }
        output[i] = sum * inv_w;
    }
}
```

3.2 OpenMP Multi-Core Shared Memory

The OpenMP implementation leverages multi-core symmetric multiprocessing (SMP) by distributing outer-loop iterations across available CPU cores:

```
void moving_average_omp(const double* input, double* output, int N, int window_size, int num_threads) {
    int k = (window_size - 1) / 2;
    double inv_w = 1.0 / window_size;
    #pragma omp parallel for num_threads(num_threads) schedule(static) default(none) \
        shared(input, output, N, k, inv_w)
    for (int i = 0; i < N; ++i) {
        double sum = 0.0;
        for (int j = -k; j <= k; ++j) {
            int idx = i + j;
            if (idx < 0) idx = 0;
            else if (idx >= N) idx = N - 1;
            sum += input[idx];
        }
        output[i] = sum * inv_w;
    }
}
```

OpenMP utilizes a static block partition where each thread receives a contiguous slice of approximately $\lfloor N/T \rfloor$ elements. Because signal elements are read-only and each thread writes strictly to disjoint output memory locations $y[i]$, no locks or atomic operations are required.

3.3 CUDA Baseline (Global Memory SIMT)

The CUDA baseline maps each output element $y[i]$ to an independent GPU thread:

- Thread index calculation: `int i = blockIdx.x * blockDim.x + threadIdx.x;`
- Execution grid: $\lceil N/256 \rceil$ thread blocks, each containing 256 threads.
- Memory access: Every thread performs W global memory loads directly from device DRAM (`d_input`). While adjacent threads access consecutive indices yielding coalesced transactions, each input element is redundantly re-fetched from high-latency global memory up to W times across neighboring threads.

3.4 CUDA Optimized (Shared-Memory Tiling with Halo Cells)

To eliminate redundant global memory access latency, the CUDA optimized implementation introduces cooperative on-chip shared-memory caching:

- Each block of $B = 256$ threads cooperatively loads its local tile of 256 elements plus $2k$ halo elements (k left neighbors and k right neighbors) into high-speed on-chip SRAM (`__shared__ double s_data[BLOCK_SIZE + 2 * MAX_K]`).
- The block synchronizes via `__syncthreads()`.
- Each thread then computes its moving-average summation strictly out of on-chip shared memory at an order-of-magnitude lower access latency.

```
+----- Shared Memory Block Tile (B + 2k) -----+
| Left Halo (k) | Interior Thread Data (B=256) | Right Halo (k) |
+-----+-----+-----+-----+-----+-----+
```

3.5 Basic MPI: 1D Domain Decomposition & Blocking Halo Exchange

In distributed-memory architectures, processes do not share an address space. The global input array of length N is decomposed across P MPI ranks. Each rank $r \in [0, P - 1]$ owns a local partition of size $N_{\text{local}} \approx N/P$.

To evaluate the moving-average stencil at local partition boundaries:

- Elements near the left boundary require up to k samples owned by rank $r - 1$.
- Elements near the right boundary require up to k samples owned by rank $r + 1$.

```
Rank r-1                      Rank r (Local Domain)                      Rank r+1
[... | Right Halo (k)] ----> [Ghost L (k) | Local Elements | Ghost R (k)] <--- [Left Halo (k) | ...]
```

In the Basic MPI implementation, processes exchange ghost cells synchronously using `MPI_Sendrecv`:

1. Rank r sends its leftmost k elements to rank $r - 1$ while receiving rank $r + 1$'s leftmost k elements into its right ghost buffer.
2. Rank r sends its rightmost k elements to rank $r + 1$ while receiving rank $r - 1$'s rightmost k elements into its left ghost buffer.
3. Ranks 0 and $P - 1$ clamp against physical array boundaries (rank 0 replicates $x[0]$; rank $P - 1$ replicates $x[N - 1]$).
4. After both `MPI_Sendrecv` calls finish, the process performs the local moving-average computation across all N_{local} elements.

Because `MPI_Sendrecv` blocks until message transfers are buffered or initiated, communication and computation are strictly serialized.

3.6 Optimized MPI: Non-Blocking Halo Exchange with Computation Overlap

The Optimized MPI implementation decouples communication from computation using non-blocking primitives (`MPI_Irecv` and `MPI_Isend`):

1. **Asynchronous Communication Post:** Rank r immediately posts non-blocking receives for its left and right ghost regions, followed by non-blocking sends of its local boundary data.
2. **Interior Computation (Window of Overlap):** While halo exchanges transfer across the communication fabric in the background, Rank r immediately computes the moving-average for its **interior domain**:

$$i_{\text{local}} \in [k, N_{\text{local}} - 1 - k]$$

These interior elements depend strictly on local data already present in memory and do not require ghost cells.

3. **Synchronization Barrier:** Rank r calls `MPI_Waitall` to ensure incoming ghost cells have arrived.
4. **Boundary Computation:** Rank r computes the remaining boundary regions:
 - Left boundary: $i_{\text{local}} \in [0, \min(k - 1, N_{\text{local}} - 1)]$
 - Right boundary: $i_{\text{local}} \in [\max(0, N_{\text{local}} - k), N_{\text{local}} - 1]$

```
Timeline: Optimized MPI Overlap Mechanism
+-----+-----+-----+-----+-----+-----+
| Post MPI_Irecv / MPI_Isend (Non-Blocking) |
+-----+-----+-----+-----+-----+-----+
| Compute Interior Domain [k, N_local - 1 - k] | MPI Communication transfers in background |
+-----+-----+-----+-----+-----+-----+
| MPI_Waitall() <-- Synchronization barrier (stalls only if communication duration > interior) |
+-----+-----+-----+-----+-----+-----+
| Compute Left & Right Boundary Domains [0, k-1] & [N_local - k, N_local - 1] |
+-----+-----+-----+-----+-----+-----+
```

4. Correctness Verification & Numerical Equivalence Methodology

Verification Standard

In scientific computing and parallel signal processing, parallelized implementations must preserve absolute fidelity to the numerical reference. For floating-point algorithms, differences can arise if associative reordering alters rounding behavior. In our moving-average filter, every parallel implementation executes identical sequential accumulation order within each window, meaning numerical divergence should be within standard machine precision limits.

Correctness was validated by comparing each parallel output array y_{parallel} against the sequential baseline y_{seq} using two metrics:

1. **Maximum Absolute Error (L_∞ Norm):**

$$\text{MaxError} = \max_{0 \leq i < N} |y_{\text{parallel}}[i] - y_{\text{seq}}[i]|$$

2. **Root-Mean-Square Error (RMSE):**

$$\text{RMSE} = \sqrt{\frac{1}{N} \sum_{i=0}^{N-1} (y_{\text{parallel}}[i] - y_{\text{seq}}[i])^2}$$

The acceptance threshold for numerical validity was set to $\varepsilon \leq 1.0 \times 10^{-4}$.

Empirical Verification Results

All six implementations were tested across Small ($N = 10,000, W = 15$), Medium ($N = 100,000, W = 31$), and Large ($N = 1,000,000, W = 63$) datasets. In addition, the MPI test suites executed dedicated boundary-condition assertions using 2, 4, 8, 12, 16, and 24 process configurations.

[!IMPORTANT]

Audited Numerical Equivalence Statement:

All parallel implementations achieved rigorous numerical equivalence to the sequential reference baseline, recording a maximum absolute error of 0.00e+00 and root-mean-square error of 0.00e+00 within floating-point tolerance (epsilon <= 1e-4) across all tested signal lengths.

Every parallel implementation passed with 100% validation across all workloads.

5. Experimental Environment & Hardware Specifications

All empirical benchmarks were conducted on a high-performance heterogeneous workstation. The experimental platform specifications are summarized below:

System Component	Hardware / Software Specification	Architectural Notes
Host Processor (CPU)	Intel Core i7-13700HX (Raptor Lake Mobile)	16 Physical Cores (8 Performance-cores + 8 Efficient-cores), 24 Logical Threads
P-Core / E-Core Clocks	P-Core Base: 2.1 GHz (Boost: 5.0 GHz) / E-Core Base: 1.5 GHz	Hybrid heterogeneous core architecture with shared L3 cache
Host Memory (RAM)	16 GB DDR5 Synchronous DRAM	High-bandwidth dual-channel host interface
Discrete Accelerator (GPU)	NVIDIA GeForce RTX 4050 Laptop GPU	6 GB GDDR6 Device Memory, Ada Lovelace Architecture (sm_89)
GPU Streaming Multiprocessors	20 SMs, 2560 CUDA Cores	Peak FP32: ~9.0 TFLOPS, Boost Clock: ~2370 MHz
Operating System	Microsoft Windows 11 Home (Build 26100, x86_64)	High-performance power profile active during runs
C++ Host Compiler	MSYS2 UCRT64 GCC 16.1.0 (g++)	Flags: -O3 -march=native -Wall -std=c++17 -fopenmp
CUDA Toolkit / Compiler	NVIDIA CUDA Toolkit 13.4 (nvcc)	Driver: 616.92, Architecture Flag: -arch=sm_89 -O3
MPI Implementation	Microsoft MPI (MS-MPI) v10.1.12498.18	Native Windows 64-bit MPI runtime and headers
Profiling & Telemetry Tools	NVIDIA Nsight Systems (nsys), Intel RAPL via PDH, NVML	Hardware energy and instruction timing collectors

6. Benchmarking & Measurement Methodology

To ensure measurement repeatability and eliminate statistical anomalies:

1. **Warmup Invocations:** Every benchmark executable performed 3 un-timed warmup passes to ensure cold-start page faults, GPU driver context initialization, and MPI daemon thread initialization did not contaminate measured durations.

2. **High-Resolution Timers:**

- Host CPU benchmarks (Sequential, OpenMP) utilized `std::chrono::high_resolution_clock`.

- CUDA benchmarks utilized hardware-synchronized CUDA Events (`cudaEventRecord`, `cudaEventSynchronize`, `cudaEventElapsedTime`) to measure pure kernel execution separate from host-to-device (H2D) and device-to-host (D2H) memory transfers.

- MPI benchmarks utilized `MPI_Wtime()` synchronized across all ranks using `MPI_Barrier(MPI_COMM_WORLD)`.

3. **Controlled Workload Configurations:**

- **Small:** $N = 10,000$ samples, $W = 15$ ($k = 7$)

- **Medium:** $N = 100,000$ samples, $W = 31$ ($k = 15$)

- **Large:** $N = 1,000,000$ samples, $W = 63$ ($k = 31$)

7. MPI Strong Scaling & Parallel Efficiency Analysis

7.1 Absolute Speedup vs. Self-Consistent Parallel Scaling

In distributed computing, scalability must be evaluated with mathematical rigor:

- **Absolute Speedup** (S_{abs}): Measures the speedup of an MPI execution of P processes relative to the standalone sequential CPU baseline:

$$S_{\text{abs}}(P) = \frac{T_{\text{seq}}}{T_{\text{mpi}}(P)}$$

- **Absolute Efficiency** (E_{abs}):

$$E_{\text{abs}}(P) = \frac{S_{\text{abs}}(P)}{P} \times 100\%$$

- **Self-Consistent Parallel Speedup** (S_{rel}): Evaluates parallel scaling against the single-process MPI baseline ($P = 1$):

$$S_{\text{rel}}(P) = \frac{T_{\text{mpi}}(1)}{T_{\text{mpi}}(P)}$$

- **Self-Consistent Parallel Efficiency** (E_{rel}):

$$E_{\text{rel}}(P) = \frac{S_{\text{rel}}(P)}{P} \times 100\%$$

[!NOTE]

Sequential Baseline and Benchmark Session Distinctions:

The sequential baseline timing recorded during the dedicated MPI scaling benchmark session was **35.027 ms** (for $N = 1,000,000$, $W = 63$). This is distinct from the **42.515 ms** sequential measurement recorded in the multi-paradigm six-way benchmark session. They represent separate benchmark invocations with normal run-to-run system and thermal variation.

Likewise, the parallel MPI execution timings recorded during the dedicated strong-scaling benchmark session (e.g., Optimized MPI $P=16$ at 5.600 ms and Basic MPI $P=16$ at 6.278 ms in Section 7.2) differ slightly from the measurements recorded during the comprehensive six-way comparison session (Optimized MPI $P=16$ at 7.584 ms and Basic MPI $P=16$ at 7.512 ms in Section 11.1). These represent separate benchmark invocations with expected run-to-run variations in operating system process launch, inter-process communication pipe initialization, and core scheduling.

[!IMPORTANT]


Audited Scaling Statement (Addressing >100% Absolute Efficiency):

In the Large workload ($N = 1,000,000$), the absolute speedup exceeds P at $P = 1, 2, 4$ (absolute efficiencies of 138.7%, 130.1%, and 110.4%). This occurs because the $P = 1$ MPI executable, compiled with `-O3` under MSYS2 UCRT64 GCC with distinct loop framing, achieved an execution time of 25.248 ms compared to the 35.027 ms of the standalone sequential benchmark. This difference does NOT represent superlinear algorithmic scaling. When evaluated self-consistently against the $P = 1$ MPI baseline (25.248 ms), scaling is strictly sublinear, demonstrating classic parallel efficiency trends.

7.2 Empirical MPI Scaling Results

The complete empirical strong-scaling measurements recorded in `pa2_mpi_scaling.csv` are presented below:

Workload	N	Window	Processes (P)	Seq Time (ms)	Basic Compute (ms)	Basic Total (ms)	Opt Compute (ms)	Opt Waitall (ms)	Opt Total (ms)	Opt Abs Speedup	Opt Self-Consistent Rel Speedup	Opt Self-Consistent Efficiency
Medium	100,000	31	1	1.487	1.091	1.809	0.725	0.000	1.277	1.16x	1.00x	100.0%
Medium	100,000	31	2	1.487	0.397	0.803	0.473	0.001	0.932	1.60x	1.37x	68.5%
Medium	100,000	31	4	1.487	0.246	0.647	0.271	0.009	0.664	2.24x	1.92x	48.1%
Medium	100,000	31	8	1.487	0.208	1.304	0.214	0.084	1.391	1.07x	0.92x	11.5%
Medium	100,000	31	12	1.487	0.137	1.616	0.156	0.155	1.154	1.29x	1.11x	9.2%
Medium	100,000	31	16	1.487	0.134	1.768	0.126	0.255	1.418	1.05x	0.90x	5.6%
Medium	100,000	31	24	1.487	0.251	5.341	0.085	0.744	2.288	0.65x	0.56x	2.3%
Large	1,000,000	63	1	35.027	19.401	25.546	19.278	0.003	25.248	1.39x	1.00x	100.0%
Large	1,000,000	63	2	35.027	9.194	13.632	9.158	0.007	13.463	2.60x	1.88x	93.8%
Large	1,000,000	63	4	35.027	5.199	9.214	4.909	0.009	7.933	4.42x	3.18x	79.6%
Large	1,000,000	63	8	35.027	2.244	5.679	2.745	0.005	5.638	6.21x	4.48x	56.0%
Large	1,000,000	63	12	35.027	3.027	6.850	4.017	0.975	12.095	2.90x	2.09x	17.4%
Large	1,000,000	63	16	35.027	2.470	6.278	2.603	0.156	5.600	6.25x	4.51x	28.2%
Large	1,000,000	63	24	35.027	1.657	7.111	1.634	2.165	7.712	4.54x	3.27x	13.6%



MPI Strong Scaling Analysis across Process Counts

7.3 Analysis of Scaling Bottlenecks Beyond $P = 8$

As observed in the scaling table:

- Linear Scaling Regime ($P \leq 8$):** For the Large dataset, performance scales steadily up to 8 processes, achieving a speedup of $6.21\times$ and an execution time of 5.638 ms. In this regime, each process computes over 125,000 samples, providing ample compute work relative to halo exchange overhead.
- Diminishing Returns & Latency Exposure ($P > 8$):** At 12, 16, and 24 processes, scaling degrades. At $P = 12$, execution time increases to 12.095 ms. At $P = 24$, execution time increases to 7.712 ms.
- Physical Core Over-subscription:** The Intel Core i7-13700HX features 16 physical cores (8 Performance-cores + 8 Efficient-cores) and 24 logical hyperthreads.

[!NOTE]

Audited Qualification on $P = 24$ Performance:


The degradation observed at $P = 24$ is consistent with increased scheduling and inter-process synchronization overhead beyond the physical-core count, alongside asymmetric execution rates between P-cores and E-cores. Because dedicated hardware performance counters for hyperthread scheduling and cache miss rates were not gathered, this phenomenon is characterized by observed inter-process synchronization stalls rather than speculative hardware cache partition claims.

8. MPI Communication Overlap & Latency Hiding Analysis

8.1 Empirical Overlap Profiling

To evaluate how effectively non-blocking communication hides halo exchange latency, the Optimized MPI implementation instrumented the exact time spent in interior computation versus `MPI_waitall` stalls. The results from `pa2_profiling_summary.csv` are tabulated below:

Processes (P)	Interior Compute (ms)	Boundary Compute (ms)	Total Compute (ms)	Waitall Stall Time (ms)	Overlap Efficiency (%)
2	0.688	0.001	0.688	0.002	99.71%
4	0.270	0.000	0.271	0.021	92.78%
8	0.235	0.001	0.236	0.003	98.74%
16	0.134	0.002	0.135	0.453	22.83%
24	0.333	0.001	0.334	0.959	25.77%



8.2 Architectural Interpretation of Overlap Dynamics

[!IMPORTANT]

Audited MPI Overlap Statement:

At moderate process counts ($P \leq 8$), non-blocking halo exchange effectively hides most measured communication overhead behind interior stencil computation. At higher process counts, the reduced partition size shortens the interior computation window, causing more communication latency to become exposed.

- Near-Perfect Hiding at $P \leq 8$:** At $P = 2$, the interior computation duration (0.688 ms) is substantially longer than the inter-process message transit time. When `MPI_waitall` is invoked, the incoming ghost messages have already arrived in the receive buffer, resulting in a negligible stall of only 0.002 ms (99.71% overlap efficiency). At $P = 8$, interior computation takes 0.235 ms while `MPI_waitall` stalls for only 0.003 ms (98.74% efficiency).
- Loss of Overlap at High Process Counts ($P \geq 16$):** At $P = 16$, the local partition length shrinks to $N_{\text{local}} \approx 62,500$. The interior computation window narrows to only 0.134 ms. Because inter-process message dispatch and kernel scheduling across 16 OS processes require hundreds of microseconds, communication cannot finish before interior computation completes. Consequently, `MPI_waitall` stalls for 0.453 ms, causing overlap efficiency to plummet to 22.83%. At $P = 24$, the stall increases to 0.959 ms (25.77% efficiency).

9. GPU Profiling & NVIDIA Nsight Systems Deep Dive

9.1 Profiling Instrumentation Setup

CUDA execution was profiled using two complementary methods:

- In-Application CUDA Events:** Measuring exact device-side kernel execution and synchronous `cudaMemcpy` transfers for H2D and D2H.
- NVIDIA Nsight Systems (`nsys`):** Capturing low-level hardware performance traces, warp execution states, and instruction issue metrics.

9.2 Nsight Systems Diagnostic Findings

The Nsight Systems traces provided deep microarchitectural insights into the baseline versus optimized CUDA kernels:

Metric / Parameter	CUDA Baseline Kernel	CUDA Optimized Kernel
Profiling Workload (N)	$N = 10,000$ samples, $W = 15$	$N = 100,000$ samples, $W = 31$
Recorded Kernel Duration	$5.23\ \mu\text{s}$ (0.00523 ms)	$80.76\ \mu\text{s}$ (0.08076 ms)
Grid Dimensions	40 Blocks $\times$ 256 Threads	391 Blocks $\times$ 256 Threads
Register Usage	16 Registers / Thread	18 Registers / Thread
Shared Memory Allocation	0 Bytes / Block	1,280 Bytes / Block
Memory Access Pattern	Uncached Global Memory Loads	Cooperative SRAM Tiling + Halo Caching

[!WARNING]

Workload Disparity Qualification:
The Nsight Systems traces were collected on different workload sizes ($N = 10,000$ for Baseline vs. $N = 100,000$ for Optimized) to capture fine-grained trace timelines without buffer overflows. Therefore, the absolute kernel timings ($5.23\ \mu\text{s}$ vs. $80.76\ \mu\text{s}$) are NOT directly comparable as a speedup ratio.

9.3 Cross-Architecture CUDA Event Comparison on Identical Workloads

For direct, side-by-side comparison on identical datasets, the synchronized CUDA Event measurements from `pa2_full_comparison.csv` must be referenced:

Workload	N	Implementation	Kernel Compute Time (ms)	Transfer Time (H2D + D2H) (ms)	End-to-End Time (ms)	Transfer Overhead Ratio (%)
Small	10,000	CUDA Baseline	0.0329	0.0955	0.1284	74.4%
Small	10,000	CUDA Optimized	0.0192	0.0515	0.0708	72.7%
Medium	100,000	CUDA Baseline	0.1030	0.1781	0.2811	63.4%
Medium	100,000	CUDA Optimized	0.1136	0.4028	0.5165	78.0%
Large	1,000,000	CUDA Baseline	1.5941	2.0708	3.6649	56.5%
Large	1,000,000	CUDA Optimized	1.5911	2.8139	4.4050	63.9%

9.4 Architectural Tradeoffs: Global vs. Shared Memory

1. **Kernel Compute Parity at Large Scale:** At $N = 1,000,000$, the pure compute times of the Baseline (1.5941 ms) and Optimized (1.5911 ms) kernels are virtually identical. On the modern Ada Lovelace architecture (`sm_89`), the large hardware L2 cache (32 MB) absorbs a significant fraction of redundant global memory fetches even in the baseline kernel.
2. **PCIe Transfer Bottleneck:** For both CUDA implementations, data transfer across the PCIe bus accounts for over 55% to 78% of total end-to-end execution time. Accelerating pure kernel computation provides diminishing returns on end-to-end latency unless data transfers are overlapped using asynchronous CUDA streams or multiple filtering stages are chained directly in GPU memory.

10. Hardware Energy Telemetry & Normalized Energy Analysis

10.1 Telemetry Methodology

Energy consumption was measured directly from hardware sensors during controlled benchmark executions on the Large dataset ($N = 1,000,000$, $W = 63$):

- **CPU Energy (Intel RAPL):** Captured via Windows Performance Data Helper (`win32pdh`) querying the Intel RAPL Energy Counter for processor package power.
- **GPU Energy (NVIDIA NVML):** Captured via native calls to `nvmiDeviceGetPowerUsage` through `nvmi.dll`, sampling GPU board power at millisecond resolution.
- **Session Duration:** Process wall-clock execution time during the energy sampling session.

10.2 Energy Normalization Methodology

Because benchmark invocations for different paradigms employed different iteration counts to achieve stable hardware power readings:

- Sequential: 5 iterations
- OpenMP ($T = 24$): 10 iterations
- Basic MPI ($P = 8$): 10 iterations
- Optimized MPI ($P = 8$): 10 iterations
- CUDA Baseline: 15 iterations
- CUDA Optimized: 15 iterations

[!CAUTION]

Strict Energy Normalization Requirement:


Raw session energy measurements reflect differing iteration counts and must NOT be compared as if they represented an identical computational task. To evaluate energy efficiency fairly, all energy values must be normalized to **Joules per filtering pass**:

$$E_{\text{pass}} = \frac{E_{\text{session}}}{N_{\text{iterations}}}$$

10.3 Empirical Energy Measurements and Energy-Delay Product (EDP)

The raw telemetry from `pa2_energy_comparison.csv` alongside the normalized per-pass energy and Energy-Delay Product ($EDP_{\text{pass}} = E_{\text{pass}} \times T_{\text{pass}}$) are detailed below:

Implementation	Hardware Telemetry Source	Iterations	Session Time (s)	Average Power (W)	Raw Session Energy (J)	Normalized Energy (E_{pass})	Filter Pass Time (T_{pass})	Per-Pass EDP ($E_{\text{pass}} \times T_{\text{pass}}$)
Sequential	Intel RAPL (CPU)	5	0.2444	57.98 W	14.169 J	2.834 J/pass	42.515 ms	0.1205 J-s
OpenMP (T=24)	Intel RAPL (CPU)	10	0.1562	70.27 W	10.977 J	1.098 J/pass	6.518 ms	0.0072 J-s
Basic MPI (P=8)	Intel RAPL (CPU)	10	0.3835	69.44 W	26.632 J	2.663 J/pass	5.679 ms	0.0151 J-s
Opt MPI (P=8)	Intel RAPL (CPU)	10	0.2959	66.57 W	19.696 J	1.970 J/pass	5.638 ms	0.0111 J-s
CUDA Baseline	NVIDIA NVML (GPU)	15	0.4196	17.97 W	7.541 J	0.503 J/pass	3.665 ms	0.0018 J-s
CUDA Optimized	NVIDIA NVML (GPU)	15	0.3675	11.29 W	4.149 J	0.277 J/pass	4.405 ms	0.0012 J-s



10.4 Quantitative Analysis of Energy Findings

[!IMPORTANT]

Audited Energy Evaluation Statement:

Hardware energy measurements using Intel RAPL for CPU package power and NVIDIA NVML for GPU board power show that CUDA Optimized had the lowest measured energy per filtering pass on this platform. Normalized to a single moving-average pass on $N=1,000,000$, CUDA Optimized consumed 0.277 J/pass, compared with 2.834 J/pass for the sequential baseline and 1.098 J/pass for OpenMP.


- GPU Energy Efficiency:** The GPU implementations demonstrated superior energy efficiency per pass. While CPU implementations operated at 58 W to 70 W package power, the GPU implementations drew only 11.3 W to 18.0 W board power. Consequently, CUDA Optimized achieved a **10.2×** reduction in energy per pass compared to the sequential CPU baseline (0.277 J vs. 2.834 J).
- MPI Energy Profile:** Basic MPI ($P = 8$) consumed 2.663 J/pass, close to the sequential CPU baseline. This occurs because spinning up 8 operating system processes maintains all CPU execution units in an active high-frequency C0 power state. Optimized MPI reduced energy to 1.970 J/pass ($\sim 26\%$ reduction over Basic MPI) by shortening overall active execution duration and eliminating blocking idle waits.
- Energy-Delay Product:** When evaluating joint performance and energy via $EDP = \text{Energy} \times \text{Delay}$, CUDA Optimized achieved the lowest EDP (0.0012 J · s), followed closely by CUDA Baseline (0.0018 J · s) and OpenMP (0.0072 J · s).

11. Comprehensive Six-Way Cross-Paradigm Performance Comparison

11.1 Full Benchmark Comparison Matrix

The complete empirical comparison across all six implementations from `pa2_full_comparison.csv` is presented below:

Implementation	Configuration	Workload	N	Window	Compute Time (ms)	Transfer / Comm (ms)	End-to-End Time (ms)	Speedup (vs. Seq)	Parallel Efficiency	Numerical Max Error	Verification Status
Sequential	1 Thread	Small	10,000	15	0.152	0.000	0.152	1.00x	100.0%	0.00e+00	REFERENCE
OpenMP	24 Threads	Small	10,000	15	0.200	0.000	0.200	0.76x	3.17%	0.00e+00	PASSED
CUDA Baseline	Block=256	Small	10,000	15	0.033	0.096	0.128	1.18x	N/A	0.00e+00	PASSED
CUDA Optimized	Shared-Mem Block=256	Small	10,000	15	0.019	0.052	0.071	2.15x	N/A	0.00e+00	PASSED
Basic MPI	16 Processes	Small	10,000	15	0.009	0.474	0.483	0.31x	1.97%	0.00e+00	PASSED
Optimized MPI	16 Processes (Overlap)	Small	10,000	15	0.008	0.117	0.649	0.23x	1.46%	0.00e+00	PASSED
Sequential	1 Thread	Medium	100,000	31	1.903	0.000	1.903	1.00x	100.0%	0.00e+00	REFERENCE
OpenMP	24 Threads	Medium	100,000	31	0.449	0.000	0.449	4.24x	17.66%	0.00e+00	PASSED
CUDA Baseline	Block=256	Medium	100,000	31	0.103	0.178	0.281	6.77x	N/A	0.00e+00	PASSED
CUDA Optimized	Shared-Mem Block=256	Medium	100,000	31	0.114	0.403	0.517	3.68x	N/A	0.00e+00	PASSED
Basic MPI	16 Processes	Medium	100,000	31	0.132	2.515	2.647	0.72x	4.49%	0.00e+00	PASSED
Optimized MPI	16 Processes (Overlap)	Medium	100,000	31	0.123	0.729	2.081	0.91x	5.72%	0.00e+00	PASSED
Sequential	1 Thread	Large	1,000,000	63	42.515	0.000	42.515	1.00x	100.0%	0.00e+00	REFERENCE
OpenMP	24 Threads	Large	1,000,000	63	6.518	0.000	6.518	6.52x	27.18%	0.00e+00	PASSED
CUDA Baseline	Block=256	Large	1,000,000	63	1.594	2.071	3.665	11.60x	N/A	0.00e+00	PASSED
CUDA Optimized	Shared-Mem Block=256	Large	1,000,000	63	1.591	2.814	4.405	9.65x	N/A	0.00e+00	PASSED
Basic MPI	16 Processes	Large	1,000,000	63	2.565	4.947	7.512	5.66x	35.37%	0.00e+00	PASSED
Optimized MPI	16 Processes (Overlap)	Large	1,000,000	63	2.583	0.319	7.584	5.61x	35.04%	0.00e+00	PASSED



11.2 Paradigm Breakdown and Performance Synthesis

1. **Small Workload** ($N = 10,000$):
2. Sequential execution requires only 0.152 ms.
3. OpenMP incurs thread team barrier overhead, yielding a slowdown (0.200 ms, $0.76\times$).
4. MPI incurs inter-process socket/shared-memory pipe setup, resulting in severe slowdowns (0.483 ms for Basic MPI, 0.649 ms for Opt MPI).
5. Only CUDA Optimized achieves a modest speedup ($2.15\times$, 0.071 ms), though data transfer accounts for over 72% of elapsed time.
6. **Medium Workload** ($N = 100,000$):
7. OpenMP scales effectively, delivering a $4.24\times$ speedup (0.449 ms).
8. CUDA Baseline achieves the highest speedup at $6.77\times$ (0.281 ms).
9. MPI remains constrained by communication overhead at 16 processes (2.647 ms for Basic MPI, 2.081 ms for Opt MPI).
10. **Large Workload** ($N = 1,000,000$):
11. **CUDA Baseline** achieves the fastest overall end-to-end execution (3.665 ms, $11.60\times$ speedup).
12. **CUDA Optimized** delivers a $9.65\times$ speedup (4.405 ms end-to-end; pure kernel compute of 1.591 ms represents a $26.7\times$ compute-only speedup).
13. **OpenMP (24 Threads)** achieves a substantial $6.52\times$ speedup (6.518 ms), operating entirely in host memory with zero interconnect transfer overhead.
14. **Optimized MPI** ($P = 16$) delivers a $5.61\times$ speedup (7.584 ms), with pure computation at 2.583 ms.

12. Architectural Bottleneck Analysis & Amdahl's Law

12.1 Memory-Bound Nature of Moving-Average Filtering

The computational intensity (operational intensity) of the moving-average filter is defined as arithmetic operations per byte of memory traffic:

$$\text{Arithmetic Operations} = W \text{ additions} + 1 \text{ multiplication} = W + 1 \text{ FLOPs}$$

$$\text{Memory Traffic (Worst Case)} = W \text{ loads} \times 8 \text{ bytes} + 1 \text{ store} \times 8 \text{ bytes} = 8(W + 1) \text{ Bytes}$$

$$\text{Operational Intensity} = \frac{W + 1}{8(W + 1)} = 0.125 \text{ FLOPs/Byte}$$

Because the operational intensity is strictly constant at 0.125 FLOPs/Byte regardless of window size W , 1D moving-average filtering is fundamentally **memory-bandwidth bound**. No architecture can exceed the throughput permitted by its memory hierarchy.

12.2 Amdahl's Law and Interconnect Bottlenecks

According to Amdahl's Law, maximum achievable speedup is bounded by the non-parallelizable fraction α :

$$S_{\max}(P) = \frac{1}{\alpha + \frac{1-\alpha}{P}}$$

1. **For GPU Computing:** The non-parallelizable sequential fraction α is dominated by PCIe data transfer time ($T_{\text{transfer}} = T_{\text{H2D}} + T_{\text{D2H}}$). Even if GPU kernel time approaches zero, speedup is strictly capped by:

$$S_{\max, \text{GPU}} \leq \frac{T_{\text{seq}}}{T_{\text{transfer}}} = \frac{42.515 \text{ ms}}{2.071 \text{ ms}} \approx 20.5 \times$$

2. **For MPI Distributed Computing:** The sequential fraction α is dominated by inter-process message serialization, OS context switching, and halo exchange latency. When partition size shrinks below a critical threshold, communication duration exceeds interior computation, bounding scaling efficiency.

13. When Parallelization Does Not Help

A critical principle of parallel and distributed computing is identifying scenarios where parallelization is counterproductive:

1. **Low Problem Dimension / Short Signal Lengths:** For $N \leq 10,000$, the overhead of spawning thread pools, initializing MPI message buffers, or launching GPU kernels exceeds the serial computation time of the entire problem.
2. **Excessive Process / Thread Concurrency:** Allocating 24 MPI processes on a single laptop workstation causes memory bus contention and inter-process scheduling overhead that degrades throughput compared to 8 processes.
3. **High Communication-to-Computation Ratios:** As process count P increases under strong scaling, partition size N/P decreases while halo size k remains constant. The ratio of halo communication to local computation scales as $\mathcal{O}(k \cdot P/N)$, inevitably leading to communication bottlenecks.

14. Reproducibility Guide & Verification Workflow

To guarantee full scientific reproducibility, all source code, benchmark runners, and plotting scripts are integrated into automated workflows.

14.1 Compilation

All six implementations compile through PowerShell and native toolchains:

```
# Build all targets (Sequential, OpenMP, CUDA, Basic MPI, Optimized MPI)
.\build.ps1 -Target all
```

Compilation binaries generated in `bin/`:

```
- bin/moving_average_seq.exe
- bin/moving_average_omp.exe
- bin/moving_average_cuda.exe
- bin/moving_average_cuda_opt.exe
- bin/moving_average_mpi.exe
- bin/moving_average_mpi_opt.exe
```

14.2 Running Benchmarks & Verifying Correctness

```
# Run correctness test suite across MPI ranks
mpiexec -n 4 bin/moving_average_mpi.exe --test
mpiexec -n 4 bin/moving_average_mpi_opt.exe --test

# Execute full automated PA2 benchmark suite
python benchmarks/run_pa2_benchmarks.py

# Regenerate plots
python scripts/plot_pa2_results.py
```

All benchmark logs and tables are written to `results/tables/`.

15. Limitations, Threats to Validity & Observations

- Hardware Cache Counter Availability:** Due to standard user-level Windows OS security permissions, direct hardware performance counters for L1/L2 cache miss rates and pipeline stall cycles could not be accessed. Architectural analyses are grounded in wall-clock timings, CUDA Events, and Nsight Systems instruction statistics.
- Single-Node Shared-Memory MPI Execution:** MS-MPI was executed on a single shared-memory node rather than a multi-node cluster with dedicated InfiniBand interconnects. Communication occurred via OS shared-memory pipes and loopback IPC.
- Dynamic Thermal Throttling:** As a mobile processor, the Intel Core i7-13700HX exhibits dynamic clock frequency scaling based on thermal headroom and power distribution between CPU and discrete GPU. Warmups and repeated trials minimized but could not entirely eliminate thermal variance.

16. Viva Voce & Oral Defense Preparation

Question 1: Why does absolute efficiency exceed 100% at $P=1, 2, 4$ in MPI scaling? Does this represent superlinear scaling?

Answer: No, it does NOT represent superlinear algorithmic scaling. The standalone sequential benchmark ran at 35.027 ms, whereas the single-process MPI executable ($P = 1$) executed in 25.248 ms due to compiler optimization variations in MSYS2 GCC -O3 loop framing and memory layout. When evaluated self-consistently against the $P = 1$ MPI baseline, efficiency is strictly sublinear: 93.8% at $P = 2$, 79.6% at $P = 4$, and 56.0% at $P = 8$.

Question 2: Why does MPI overlap efficiency drop from 98.7% at $P=8$ to 22.8% at $P=16$?

Answer: At $P = 8$, each process computes over 125,000 samples. The interior computation window (0.235 ms) is long enough to completely hide non-blocking halo exchange latency, resulting in an `MPI_Waitall` stall of only 0.003 ms. At $P = 16$, the partition shrinks to 62,500 samples, narrowing interior compute to 0.134 ms. Because inter-process message dispatch and kernel scheduling require more time, communication is exposed, forcing `MPI_Waitall` to stall for 0.453 ms.

Question 3: How was energy normalized across implementations with different iteration counts?

Answer: Raw session energy reflects differing repetition counts (5 iterations for Sequential, 10 for OpenMP and MPI, 15 for CUDA) chosen to obtain stable hardware sensor readings. Energy was rigorously normalized to Joules per filtering pass: $E_{\text{pass}} = E_{\text{session}} / N_{\text{iterations}}$. This revealed that CUDA Optimized consumed 0.277 J/pass, compared to 1.098 J/pass for OpenMP and 2.834 J/pass for Sequential.

Question 4: Why do CUDA Baseline and CUDA Optimized exhibit nearly identical kernel runtimes on Large workloads?

Answer: The NVIDIA Ada Lovelace RTX 4050 GPU features a 32 MB hardware L2 cache. At $N = 1,000,000$, the input array fits within L2 cache, allowing the baseline kernel's redundant fetches to hit cache rather than DRAM. Consequently, the explicit shared-memory caching in the optimized kernel provides minimal additional latency reduction.

17. Conclusion

Programming Assignment 2 provided an extensive empirical exploration across six computational paradigms for 1D moving-average signal filtering. Key conclusions include:

- Numerical Invariance:** All six implementations maintained rigorous numerical equivalence to the sequential reference baseline, achieving Max Absolute Error = 0.00e+00 across all signal sizes.
- Computational Sweet Spots:**
 - For Small datasets ($N \leq 10,000$), Sequential execution is fastest due to zero launch and synchronization overhead.
 - For Large datasets ($N = 1,000,000$), **CUDA Baseline** achieved the highest end-to-end speedup (**11.60×**), while **OpenMP** delivered the strongest host-only performance (**6.52×**) with zero data transfer latency.
- Latency Hiding in Distributed Stencils:** Non-blocking MPI communication (`MPI_Isend` / `MPI_Irecv`) successfully hides over 98% of halo exchange latency up to 8 processes, demonstrating the viability of domain decomposition for stencil filtering.
- Energy Efficiency:** Accelerated computing on GPUs demonstrated a **10.2×** reduction in energy per filtering pass (0.277 J/pass vs. 2.834 J/pass), confirming that architectural throughput directly enhances computational energy efficiency.

18. Contributions & Team Responsibilities

Team Member	Roll Number	Key Responsibilities & Module Leadership
Saksham Singh	2024BCS0070	Project Lead; MPI domain decomposition & halo exchange design; Optimized MPI non-blocking overlap implementation; Hardware energy telemetry engine (RAPL & NVML); Final report synthesis & mathematical modeling.
Daksh Singh	2024BCS0042	Basic MPI implementation; Correctness verification suite and boundary condition testing; Nsight Systems profiling integration; Data collection across OpenMP and MPI scaling benchmarks.
Anmol Pipara	2024BCS0014	Benchmark automation scripting; Data table formatting and traceability auditing; Matplotlib plotting pipeline for PA2 scaling, overlap, and energy figures; Reproducibility verification.